

```
In [2]: import pandas as pd
```

```
In [3]: import numpy as np
```

```
In [4]: df=pd.read_csv(r'C:\Users\Rutu\Documents\PDF\housingData.csv')
```

```
In [5]: df.head()
```

```
Out[5]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	NaN

```
In [6]: df.tail()
```

```
Out[6]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	NaN
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	9.08
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	6.48
505	0.04741	0.0	11.93	0.0	0.573	6.030	NaN	2.5050	1	273	21.0	396.90	7.88

```
In [7]: df.describe()
```

```
Out[7]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE
count	486.000000	486.000000	486.000000	486.000000	506.000000	506.000000	486.000000
mean	3.611874	11.211934	11.083992	0.069959	0.554695	6.284634	68.518519
std	8.720192	23.388876	6.835896	0.255340	0.115878	0.702617	27.999513
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000
25%	0.081900	0.000000	5.190000	0.000000	0.449000	5.885500	45.175000
50%	0.253715	0.000000	9.690000	0.000000	0.538000	6.208500	76.800000
75%	3.560263	12.500000	18.100000	0.000000	0.624000	6.623500	93.975000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   CRIM        486 non-null    float64
1   ZN          486 non-null    float64
2   INDUS       486 non-null    float64
3   CHAS        486 non-null    float64
4   NOX         506 non-null    float64
5   RM          506 non-null    float64
6   AGE         486 non-null    float64
7   DIS         506 non-null    float64
8   RAD         506 non-null    int64   
9   TAX         506 non-null    int64   
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       486 non-null    float64
13  MEDV        506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

```
In [9]: df.isnull()
```

```
Out[9]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	False	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False	True
...
501	False	False	False	False	False	False	False	False	False	False	False	False	True
502	False	False	False	False	False	False	False	False	False	False	False	False	False
503	False	False	False	False	False	False	False	False	False	False	False	False	False
504	False	False	False	False	False	False	False	False	False	False	False	False	False
505	False	False	False	False	False	False	True	False	False	False	False	False	False

506 rows × 14 columns



```
In [10]: df.columns
```

```
Out[10]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',  
               'PTRATIO', 'B', 'LSTAT', 'MEDV'],  
              dtype='object')
```

```
In [11]: df.isnull().sum()
```

```
Out[11]: CRIM      20  
         ZN        20  
         INDUS    20  
         CHAS     20  
         NOX       0  
         RM        0  
         AGE      20  
         DIS       0  
         RAD       0  
         TAX       0  
         PTRATIO   0  
         B         0  
         LSTAT    20  
         MEDV      0  
         dtype: int64
```

```
In [12]: df.shape
```

```
Out[12]: (506, 14)
```

```
In [13]: df.dropna(inplace=True)
```

```
In [14]: df.isnull().sum()
```

```
Out[14]: CRIM      0  
         ZN        0  
         INDUS    0  
         CHAS     0  
         NOX       0  
         RM        0  
         AGE      0  
         DIS       0  
         RAD       0  
         TAX       0  
         PTRATIO   0  
         B         0  
         LSTAT    0  
         MEDV      0  
         dtype: int64
```

```
In [15]: df.value_counts()
```

```
Out[15]: CRIM      ZN      INDUS  CHAS  NOX      RM      AGE      DIS      RAD  TAX  PTRATIO  B
LSTAT  MEDV
0.00632  18.0  2.31  0.0  0.538  6.575  65.2  4.0900  1  296  15.3  39
6.90  4.98  24.0  1
1.34284  0.0  19.58  0.0  0.605  6.066  100.0  1.7573  5  403  14.7  35
3.89  6.43  24.3  1
1.25179  0.0  8.14  0.0  0.538  5.570  98.1  3.7979  4  307  21.0  37
6.57  21.02  13.6  1
1.23247  0.0  8.14  0.0  0.538  6.142  91.7  3.9769  4  307  21.0  39
6.90  18.72  15.2  1
1.20742  0.0  19.58  0.0  0.605  5.875  94.6  2.4259  5  403  14.7  29
2.29  14.43  17.4  1

..
0.11460  20.0  6.96  0.0  0.464  6.538  58.7  3.9175  3  223  18.6  39
4.96  7.73  24.4  1
0.11432  0.0  8.56  0.0  0.520  6.781  71.3  2.8561  5  384  20.9  39
5.58  7.67  26.5  1
0.11132  0.0  27.74  0.0  0.609  5.983  83.5  2.1099  4  711  20.1  39
6.90  13.35  20.1  1
0.11069  0.0  13.89  1.0  0.550  5.951  93.8  2.8893  5  276  16.4  39
6.90  17.92  21.5  1
88.97620  0.0  18.10  0.0  0.671  6.968  91.9  1.4165  24  666  20.2  39
6.90  17.21  10.4  1
Length: 394, dtype: int64
```

```
In [16]: df[['ZN']].value_counts()
```

```
Out[16]: ZN
0.0      291
20.0      17
80.0      10
25.0      10
22.0       9
12.5       8
90.0       5
95.0       4
45.0       4
40.0       4
33.0       3
30.0       3
60.0       3
70.0       3
75.0       3
35.0       2
28.0       2
52.5       2
55.0       2
21.0       2
82.5       2
85.0       2
18.0       1
17.5       1
100.0      1
dtype: int64
```

```
In [17]: df[['CHAS']].value_counts()
```

```
Out[17]: CHAS
0.0      367
1.0       27
dtype: int64
```

```
In [18]: df.dtypes
```

```
Out[18]: CRIM      float64
ZN          float64
INDUS      float64
CHAS       float64
NOX        float64
RM         float64
AGE        float64
DIS        float64
RAD         int64
TAX         int64
PTRATIO    float64
B          float64
LSTAT      float64
MEDV      float64
dtype: object
```

```
In [19]: y=df['CRIM']
```

```
In [20]: y
```

```
Out[20]: 0      0.00632
1      0.02731
2      0.02729
3      0.03237
5      0.02985
...
499    0.17783
500    0.22438
502    0.04527
503    0.06076
504    0.10959
Name: CRIM, Length: 394, dtype: float64
```

```
In [21]: x=df[['ZN','INDUS','CHAS','NOX','RM','AGE','DIS','RAD','TAX','PTRATIO','B','LSTAT','MEDV']]
```

```
In [22]: x
```

```
Out[22]:
```

	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
0	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
5	0.0	2.18	0.0	0.458	6.430	58.7	6.0622	3	222	18.7	394.12	5.21	28.7
...
499	0.0	9.69	0.0	0.585	5.569	73.5	2.3999	6	391	19.2	395.77	15.10	17.5
500	0.0	9.69	0.0	0.585	6.027	79.7	2.4982	6	391	19.2	396.90	14.33	16.8
502	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	9.08	20.6
503	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64	23.9
504	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	6.48	22.0

394 rows × 13 columns

```
In [23]: x.shape
```

```
Out[23]: (394, 13)
```

```
In [24]: from sklearn.model_selection import train_test_split
```

```
In [25]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.4,random_state=2)
```

```
In [26]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
Out[26]: ((236, 13), (158, 13), (236,), (158,))
```

```
In [27]: from sklearn.linear_model import LinearRegression
```

```
In [28]: lr=LinearRegression()
```

```
In [29]: lr.fit(x_train,y_train)
```

```
Out[29]: LinearRegression()
```

```
In [30]: y_pred=lr.predict(x_test)
```

```
In [31]: y_pred.shape
```

```
Out[31]: (158,)
```

```
In [32]: y_pred
```

```
Out[32]: array([12.935922 ,  1.03139699, -4.21746078, 15.37065881,  0.89075298,
                3.03311676, 16.22803816,  0.3028599 ,  4.30871359,  0.79750878,
                6.19982726,  1.25171726,  0.03758206,  1.90877592, -1.41151168,
               -0.45842888, -1.61532464,  1.29981509, -2.70302105,  1.08515213,
               17.83363067,  4.59686925,  2.1008869 ,  1.17613716, 14.48983654,
                5.03153761,  2.02889521, -0.16995535, 20.7697108 , -1.89013628,
                1.59905481, -3.82478187, -2.48892578, -2.3897628 , -0.45922083,
               -0.50957521, -2.01987301,  2.69480245, -1.43813057,  3.48991821,
                2.18396115,  0.78986788, 15.34920081,  7.71650902,  2.9751829 ,
               -0.54515732,  1.72443869, 16.73994458,  0.56656552, -1.28724361,
               18.04844372,  4.91251065, 16.99132212, -0.58645378, 15.54398917,
               -2.68280868,  0.53872057, 15.77861889,  1.00122542, -0.77916292,
               21.30193144, -0.64593524, -0.0303628 ,  3.22270199,  2.71050247,
               -0.30599732, -2.34637883,  1.10346922, -0.78161217,  1.69111522,
                2.47521546, -2.68191181,  2.31284667,  2.89030912,  4.20696855,
                0.08396164, 17.61903684, -7.52871885, -4.54305258, -1.33888227,
               -1.96025224,  1.80855902,  3.95622385, 17.93841206, -0.96922524,
                3.08395508, -2.27614313, -1.90571445,  1.26285226, 16.8538469 ,
               12.22256521,  0.42279442, 12.02868615, 13.96345377,  5.6734962 ,
               19.61189164, -3.83248688, 12.15255625, -5.03561045, -0.16179549,
               -2.11891663, -8.51705207,  2.50259354,  8.21403607,  2.67824565,
               -1.96774668,  0.72223693, 16.24830325, -3.9002718 , -0.36049838,
                0.21330248,  2.91480528, -3.51148999, -1.54047874,  2.75287094,
               17.5991266 ,  2.1164853 , -0.16043837, -1.77787688, 23.36861704,
               -4.38321192,  1.79112468, 14.9083135 ,  7.27267044, 21.59840775,
                0.64609847,  0.42677874,  0.3841648 , 17.25855024, 10.61349429,
                3.7032961 , 17.62947639, -3.42767243,  0.78833139, 15.82313715,
                0.42374552, 14.11204082, -0.19093817,  0.1038211 ,  1.42919613,
                0.1359542 , -0.29655227,  2.50152764,  1.36549247, 13.78581421,
               13.37325621, 16.95424315, -0.4363879 , -0.68717684,  2.33676091,
                2.74588406, 20.50932388, -0.74607058, -2.58931479, 13.83173767,
                1.43583166,  3.2355107 , -6.68796363])
```

```
In [33]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
In [34]: mean_squared_error(y_test, y_pred)
```

```
Out[34]: 23.57799042915025
```

```
In [35]: mean_absolute_error(y_test, y_pred)
```


```
Out[35]: 3.330071652883863
```

```
In [36]: r2_score(y_test, y_pred)
```

```
Out[36]: 0.3607529980840608
```




```
In [46]: import matplotlib.pyplot as plt
plt.scatter(y_test,y_pred)
plt.xlabel("price")
plt.ylabel("Pred_Price")
plt.show()
```



```
In [38]: import seaborn as sns
```

```
In [40]: correlation_matrix = df.corr().round(2)
# annot = True to print the values inside the square
sns.heatmap(data=correlation_matrix, annot=True)
```



Out[40]: <AxesSubplot:>



```
In [43]: plt.figure(figsize=(20, 5))

features = ['LSTAT', 'RM']
target = df['MEDV']

for i, col in enumerate(features):
    plt.subplot(1, len(features), i+1)
    x = df[col]
    y = target
    plt.scatter(x, y, marker='o')
    plt.title(col)
    plt.xlabel(col)
    plt.ylabel('MEDV')
```



```
In [44]: sns.set(rc={'figure.figsize':(11.7,8.27)})
sns.distplot(df['MEDV'], bins=30)
plt.show()
```

C:\Users\Rutu\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

